

1.Linear Regression

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import linear_model

df = pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\pizza.csv")


plt.xlabel('Size')
plt.ylabel('Price')
plt.scatter(df['size'], df['price'], color='red', marker='+')

reg = linear_model.LinearRegression()
reg.fit(df[['size']], df['price'])

plt.plot(df['size'], reg.predict(df[['size']]), color='blue')
plt.show()

predicted_price = reg.predict(np.array([[15]]))
print("Predicted price for size 15 pizza:", predicted_price[0])
```

Output:



Predicted price for size 15 pizza: 22.4

2. Logistic Regression

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

df = pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\insurance.csv") #
Removed space in filename


plt.xlabel('Age')
plt.ylabel('Insurance')
plt.scatter(df['age'], df['insurance'], color='red', marker='+')
plt.show()

x_train, x_test, y_train, y_test = train_test_split(df[['age']], df['insurance'], test_size=0.2)

model = LogisticRegression()
model.fit(x_train, y_train)

inputdata = np.array([45, 23, 56, 66, 77, 88, 99, 12])
inputdata = inputdata.reshape(-1, 1)
predictions = model.predict(inputdata)
print("Predictions:", predictions)
```

Output:



Predictions: [0 0 1 1 1 1 1 0]

3. ID3 Algorithm

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

data = pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\iris.csv")

column_names = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'class']
data.columns = column_names

print(data.isnull().sum())

X = data.drop('class', axis=1)
y = data['class']

print(f"Features shape: {X.shape}")
print(f"Target shape: {y.shape}")

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
print(f"Training set size: {X_train.shape[0]}")
print(f"Test set size: {X_test.shape[0]}")

model = DecisionTreeClassifier(criterion='entropy', random_state=42)
model.fit(X_train, y_train)

train_accuracy = model.score(X_train, y_train)
test_accuracy = model.score(X_test, y_test)

print(f"Train Accuracy: {train_accuracy:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

new_sample = [1.3, 1.3, 4.5, 8.0]
print(f"Sample length: {len(new_sample)}")

prediction = model.predict([new_sample])
print(f"Predicted Class: {prediction[0]}")
```

Output:

```
sepal_length    0
sepal_width     0
```

```
petal_length  0
petal_width   0
class         0
dtype: int64

Features shape: (150, 4)
Target shape: (150,)
Training set size: 120
Test set size: 30
Train Accuracy: 1.0000
Test Accuracy: 1.0000
Sample length: 4
Predicted Class: Virginica
```

4. Navie Bayes

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
data = pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\iris.csv")
x = data.drop(["variety"], axis=1)
y = data["variety"]
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(x)
X_train, X_test, y_train, y_test = train_test_split(x_scaled, y, test_size=0.2,
random_state=42)
gnb = GaussianNB()
gnb.fit(X_train, y_train)
y_pred = gnb.predict(X_test)
print(y_pred)
```

Output:

```
['Versicolor' 'Setosa' 'Virginica' 'Versicolor' 'Versicolor' 'Setosa'
 'Versicolor' 'Virginica' 'Versicolor' 'Versicolor' 'Virginica' 'Setosa'
 'Setosa' 'Setosa' 'Setosa' 'Versicolor' 'Virginica' 'Versicolor']
```

```
'Versicolor' 'Virginica' 'Setosa' 'Virginica' 'Setosa' 'Virginica'
'Virginica' 'Virginica' 'Virginica' 'Virginica' 'Setosa' 'Setosa']
```

5. KNN

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import LabelEncoder
df=pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\knn.csv")
x=df[['weight','height']].values
y=df['class'].values
le=LabelEncoder()
y_encoded=le.fit_transform(y)
knn=KNeighborsClassifier(n_neighbors=3)
knn.fit(x,y_encoded)
sample=[[23,56]] #Example height and weight
prediction=knn.predict(sample)
predicted_table=le.inverse_transform(prediction)[0]
print(f'Predicted status for {sample[0]} (height,weight): {predicted_table}')
```

Output:

Predicted status for [23, 56] (height,weight):underweight

6. K Means

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

df = pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\kmeans.csv")

x= df[['Height', 'Weight']]
Kmean = KMeans(n_clusters=3, random_state=30)
Kmean.fit(x)

x['cluster'] = Kmean.predict(x)

print(x)
```

Output:

Height Weight cluster

| | | | |
|---|-----|----|---|
| 0 | 185 | 72 | 0 |
| 1 | 170 | 56 | 1 |
| 2 | 168 | 60 | 1 |
| 3 | 179 | 68 | 0 |
| 4 | 182 | 72 | 0 |
| 5 | 188 | 77 | 2 |
| 6 | 189 | 71 | 2 |

7. Hierarchical Clustering

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from scipy.cluster.hierarchy import dendrogram, linkage
from sklearn.cluster import AgglomerativeClustering
from sklearn.decomposition import PCA

data = pd.read_csv("C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data
Set\\hierarchial.csv")
print(data)

x = data['x']
y = data['y']
n = range(1, 8)

fig, ax = plt.subplots()
ax.scatter(x, y, marker='*', c='red', alpha=0.5)

linked = linkage(data[['x', 'y']], 'single')

plt.figure(figsize=(10, 7))
dendrogram(linked, orientation='top', distance_sort='descending', show_leaf_counts=True)
plt.xlabel('Samples')
plt.ylabel('Euclidean Distances')
plt.show()

model = AgglomerativeClustering(n_clusters=5, metric='euclidean', linkage='single')
model.fit(data[['x', 'y']])
```

```

x = data['x']
y = data['y']
n = range(1, 8)

fig, ax = plt.subplots()
ax.scatter(x, y, c=model.labels_, cmap='rainbow')
plt.grid()
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Hierarchical Clustering')
plt.show()


```



Output:

```

Unnamed: 0    x    y
0      p1  0.40  0.53
1      p2  0.22  0.38
2      p3  0.35  0.32
3      p4  0.26  0.19
4      p5  0.08  0.41
5      p6  0.45  0.30

```





8. Support Vector Machine

```
import matplotlib.pyplot as plt
import numpy as np
from numpy.ma.extras import unique
from sklearn.svm import SVC
x=np.array([(1,0.5),(1,1),(1,-0.5),(-0.5,0.5),(0.5,0.5),(2,0),(4,0),(4.5,1),(5,-1),(5.5,0)])
y=np.array([-1,-1,-1,-1,-1,-1,1,1,1,1])

x_coords=x[:,0]
y_coords=x[:,1]
plt.figure(figsize=(8,6))
plt.scatter(x_coords[y== -1],y_coords[y== -1],color='blue',label='Class 0')
plt.scatter(x_coords[y== 1],y_coords[y== 1],color='red',label='Class 1')

clf=SVC(kernel='linear',C=1.0)
clf.fit(x,y)

support_vectors=clf.support_vectors_
for sv in support_vectors:
    plt.axvline(x=sv[0],color='red',linestyle='--',alpha=0.7,label='Support Vector Line')

handles, labels=plt.gca().get_legend_handles_labels()
unique=dict(zip(labels,handles))
plt.legend(unique.values(),unique.keys())

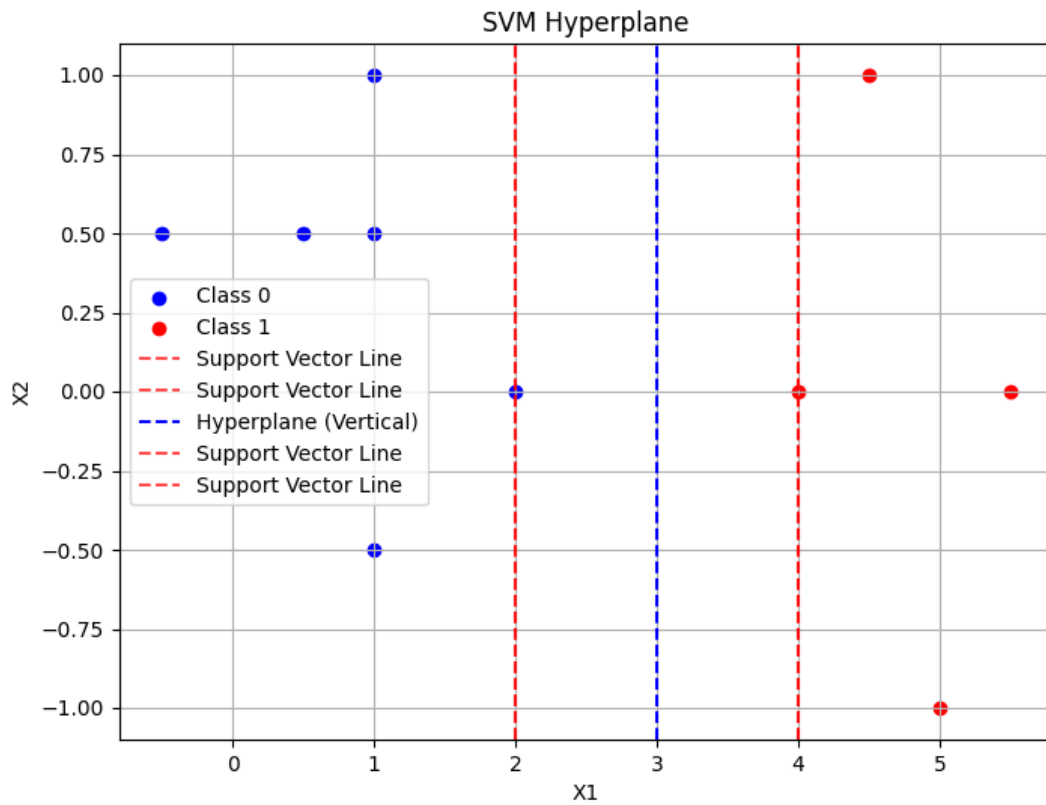
w=clf.coef_[0]
b=clf.intercept_[0]
if np.isclose(w[1],0):
    x_hyperplane=-b/w[0]
    plt.axvline(x=x_hyperplane,color='blue',linestyle='--',label='Hyperplane (Vertical)')
else:
    x_vals=np.linspace(-1,7,200)
    y_vals=(w[0]*x_vals+b)/w[1]
    plt.plot(x_vals,y_vals,'k--',label='Hyperplane')

plt.scatter(clf.support_vectors_[ :,0],clf.support_vectors_[ :,1],s=120,facecolors='none',edgecolors='blue',label='Support Vectors')
support_vectors=clf.support_vectors_
for sv in support_vectors:
    plt.axvline(x=sv[0],color='red',linestyle='--',alpha=0.7,label='Support Vector Line')

plt.xlabel("X1")
plt.ylabel("X2")
plt.title("SVM Hyperplane")
plt.legend()
```

```
plt.grid(True)
plt.show()
```

Output:



9. Random Forest

```
import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.cluster import AgglomerativeClustering
```

```
x_real = np.array([
    [1.0,2.0],
    [1.1,2.1],
    [8.0,9.0]
])
```

```
n_samples, n_features = x_real.shape
```

```

x_synthetic=np.zeros_like(x_real)
for i in range(n_features):
    x_synthetic[:,i]=np.random.permutation(x_real[:,i])

x_combined=np.vstack([x_real,x_synthetic])
y_labels=np.array([1]*n_samples+[0]*n_samples)

rf=RandomForestClassifier(n_estimators=10,max_depth=3,random_state=42)
rf.fit(x_combined,y_labels)

leaf_indices=rf.apply(x_real)

proximity=np.zeros((n_samples,n_samples))
for tree_leaf in leaf_indices.T:
    for i in range(n_samples):
        for j in range(n_samples):
            if tree_leaf[i]==tree_leaf[j]:
                proximity[i,j]+=1

proximity/=rf.n_estimators

distance=1-proximity

clustering=AgglomerativeClustering(n_clusters=2,metric='precomputed',linkage='average')
labels=clustering.fit_predict(distance)

print("Real Data Points:")
print(x_real)

print("\nProximity Matrix:")
print(np.round(proximity,2))

print("\nDistance Matrix:")
print(np.round(distance,2))

print("\nCluster Labels:")
print(labels)

```

Output:

Real Data Points:

```
[[1. 2. ]
```

```
[1.1 2.1]
```

```
[8. 9. ]]
```

Proximity Matrix:

```
[[1. 0.4 0. ]  
 [0.4 1. 0. ]  
 [0. 0. 1. ]]
```

Distance Matrix:

```
[[0. 0.6 1. ]  
 [0.6 0. 1. ]  
 [1. 1. 0. ]]
```

Cluster Labels:

```
[0 0 1]
```

10. LOF

```
import numpy as np  
from sklearn.neighbors import NearestNeighbors
```

```
data=np.array([1,2,2.5,3,3.5,10]).reshape(-1,1)  
k=2
```

```
nbrs=NearestNeighbors(n_neighbors=k+1)  
nbrs.fit(data)  
distances, indices = nbrs.kneighbors(data)
```

```
lrd=[]  
lof=[]
```

```
for i in range(len(data)):  
    reach_dists=[]  
    for j in range(1,k+1):  
        neighbor_idx=indices[i][j]  
        neighbor_k_dist=distances[neighbor_idx][k]  
        actual_dist=distances[i][j]  
        reach_dist=max(neighbor_k_dist,actual_dist)  
        reach_dists.append(reach_dist)  
    lrd_i=1/(np.mean(reach_dists))  
    lrd.append(lrd_i)
```

```

for i in range(len(data)):
    lrd_ratios=[]
    for j in range(1,k+1):
        neighbor_idx=indices[i][j]
        lrd_ratios.append(lrd[neighbor_idx]/lrd[i])
    lof_i=np.mean(lrd_ratios)
    lof.append(lof_i)

print(f'{"Point":>6} | {"LRD":>8} | {"LOF":>8} | Outlier?')
print("-"*36)
for i in range(len(data)):
    outlier='YES' if lof[i]>1.5 else 'NO'
    print(f'{"data[i][0]:>6} | {"lrd[i]:>2.2f} | {"lof[i]:>8.2f} | {outlier}')
```

Output:

```
Point |    LRD |    LOF | Outlier?
```

```
-----
```

```

1.0 | 0.80 | 1.46 | NO
2.0 | 1.00 | 1.07 | NO
2.5 | 1.33 | 0.88 | NO
3.0 | 1.33 | 1.00 | NO
3.5 | 1.33 | 1.00 | NO
10.0 | 0.15 | 9.00 | YES
```

11. Moving Average

```

import matplotlib.pyplot as plt

input=[15,16,24,28,18,21,17,22,25,16,24,15,23,22]
def moving_avg(input, window_size):
    result=[]
    moving_sum=sum(input[:window_size])
    result.append(moving_sum/window_size)
    for i in range(len(input)-window_size):
        moving_sum=moving_sum+(input[i+window_size]-input[i])
        result.append(moving_sum/window_size)
    return result
window_size=3
print(moving_avg(input,window_size))
```

```

ma=moving_avg(input>window_size)

x_ma=list(range(window_size-1,len(input)))
plt.figure(figsize=(10,5))
plt.plot(input,label='Original Data',marker='o')

plt.plot(x_ma,ma,label=f'{window_size}-Point Moving
Average',color='red',marker='o',linestyle='--')
plt.title('Moving Average Plot')
plt.xlabel('Time Index')
plt.ylabel('Value')
plt.grid(True)
plt.tight_layout()
plt.show()

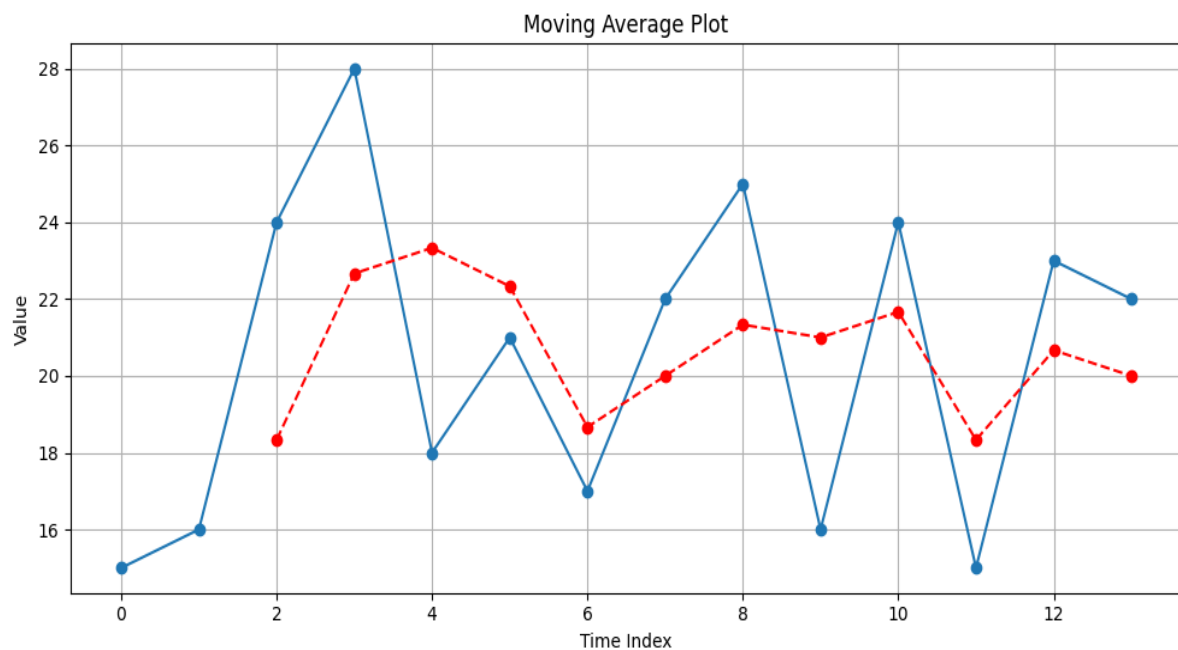
```

Output:

```

[18.333333333333332, 22.666666666666668, 23.333333333333332, 22.333333333333332,
18.666666666666668, 20.0, 21.333333333333332, 21.0, 21.666666666666668,
18.333333333333332, 20.666666666666668, 20.0]

```



12. ARIMA

```
import pandas as pd
from statsmodels.tsa.arima.model import ARIMA
import matplotlib.pyplot as plt
file_path="C:\\Users\\ASUS\\OneDrive\\Documents\\ML Data Set\\airline_passengers.csv"
df=pd.read_csv(file_path,parse_dates=['Month'],index_col='Month')
df.index.freq='MS'

model=ARIMA(df['Passengers'],order=(2,1,2))
model_fit=model.fit()
forecast=model_fit.forecast(steps=1).iloc[0]
print(f"Forecasted next value using ARIMA: {forecast:.2f}")

plt.plot(df['Passengers'],label='Original Data')
plt.plot(model_fit.fittedvalues,label='Fitted values')
plt.legend()
plt.title('ARIMA Forecasting on Airline Passengers')
plt.show()
```

Output:

Forecasted next value using ARIMA:439.85

